The Complete Record of Making a 39-Minute Psychological Thriller — 100% AI-Generated, Solo

Koki Yoshikawa — Kyoto University, Faculty of Science

About the Film

In the summer of 2025, I completed a 39-minute feature film in about one month. Cast, locations, and cameras were all replaced by AI — produced entirely from a single laptop at home.

Runtime	**39 min (feature-length)**
Video & Audio	**100% AI-generated**
Budget	**¥84,261 / ~$581**
Crew	**1 person (planning through editing)**
Hardware	**NEC LAVIE / Core i3 / 8 GB RAM / No GPU**

The genre is psychological suspense. A young journalist investigating a terrorist incident from a small apartment gradually finds herself consumed by the darkness of the human subconscious. It spans 17 sequences, 8 characters, 12 BGM tracks, and 24 sound effects.

The concept dates to 2017, when I watched Alex Garland's *Ex Machina* in high school. After eight years, I started production in the summer of 2025 as AI technology matured.

Left: Daily life scene — Right: Hallucination scene — All photorealistic footage generated by AI

Production Setup & Cost

PC Specs

I used an ordinary laptop bought when I entered university. No dedicated GPU — all production ran on cloud-based AI services.

Item	Specification
Model	NEC LAVIE
CPU	Intel Core i3-10110U (10th gen)
RAM	8 GB
GPU	Intel UHD Graphics (integrated)
Storage	~240 GB SSD

Tools Used (16 Total)

Stage	Tool	Purpose
Image Gen	Midjourney V7	Base image generation
	Flux1.Kontext	Photorealistic conversion
	Gemini 2.5 Flash image	31 expression variants
	Flux1.dev (Replicate)	LoRA training
	Imagen 3 / Imagen 4	Non-protagonist images
	ChatGPT (DALL-E)	Non-protagonist images
	Topaz Photo AI 4	4K upscaling
Video Gen	Runway Gen-3/Gen-4	Lip-sync + static scenes
	Pixverse V5	High-motion scenes
Voice Gen	Hailuo Audio	Dialogue (emotional delivery)
	Gemini 2.5 Pro TTS	Narration
Music & SFX	SUNO	BGM (strong prompt comprehension)
	Udio	BGM (excels at eerie tracks)
	ElevenLabs	BGM + SFX (precise duration control)
Editing	Adobe Premiere Pro	Final editing & assembly
	Microsoft Clipchamp	10-second audio trimming (free)

The key was matching each tool to its strength: Runway Gen-4 for static scenes, Pixverse V5 for high-motion. For music, I tested the same prompt across three tools and picked the best.

Cost Breakdown

Total: ¥84,261 / ~$581. Trial-only services (e.g. KlingAI) excluded.

Category	Service	Cost
Image Gen	Midjourney	¥8,290 / ~$57
	Black Forest Labs (Flux)	¥3,054 / ~$21
	Replicate (LoRA × 5 runs)	¥6,632 / ~$46
Upscaling	Topaz Photo AI 4	¥5,394 / ~$37
Video Gen	Runway	¥22,947 / ~$158
	Pixverse	¥6,093 / ~$42
Voice & Text	OpenAI (ChatGPT)	¥10,076 / ~$69
Music & SFX	SUNO / Udio	¥3,028 / ~$21
	ElevenLabs	¥2,698 / ~$19
Image Gen	Gemini	¥5,800 / ~$40
Voice Gen	Hailuo AI	¥1,504 / ~$10
Editing	Adobe Premiere Pro	¥6,560 / ~$45
	Total	**¥84,261 / ~$581**

The largest cost was video generation (Runway + Pixverse: ~¥29,000 / ~$200) — generating 39 minutes in 10-second clips consumed heavy credits. Voice was cheap: Hailuo Audio + Gemini TTS covered all dialogue.

Production Workflow

Production splits into three tracks: image, video, and audio generation.

Image Pipeline (Protagonist Scenes)

Midjourney V7 ① Base image generation
↓
Flux1.Kontext ② Photorealistic conversion
↓
Gemini Flash ③ Generate 31 expressions
↓
Replicate Flux1.dev ④ LoRA training
↓
LoRA + Prompt ⑤ Scene images 1280×720
↓
Topaz Photo AI 4 ⑥ Upscale to 4K (3840×2160)

Image Pipeline (Non-Protagonist Scenes)

> **Imagen 3 / 4** — Faithful to prompts
> **Gemini Flash image** — Backgrounds & props
> **ChatGPT (DALL-E)** — Supplementary generation

→ All upscaled to 4K via Topaz Photo AI 4

Video Pipeline

4K Image Upscaled
↓
image-to-video Image to video
↓
Premiere Pro Final editing & assembly

> **Runway Gen-3 / Gen-4** — Dialogue + static scenes (lip-sync, 4K output)
> **Pixverse V5** — High-motion scenes (compensates for Gen-4 limitations)

Audio Pipeline

> **Hailuo Audio** — Dialogue (natural Japanese + emotional nuance)
> **Gemini 2.5 Pro TTS** — Narration (reading-optimized, free)
> **ElevenLabs** — Sound effects (high-quality from prompts)

> **BGM (compared 3 tools → picked the best):**
> **SUNO** — Strong prompt comprehension / weak on dark tones
> **Udio** — Excels at eerie tracks / imprecise duration
> **ElevenLabs** — Second-level duration control / higher credit cost

The most labor-intensive step was ④ LoRA training — five iterations to find optimal parameters. For video, the central challenge was building 39 minutes from 10-second clips (Section 6).

Full prompts, parameter settings, and usage guides for each tool are in the paid course.

Challenges and Solutions

Completing a 39-minute film meant hitting numerous technical walls. Here are the three most critical.

Challenge 1: The Face Changed Every Time

What happened

AI image generators produce a different face every run, even with the same prompt. A film requires consistent identity — this was a fundamental blocker.

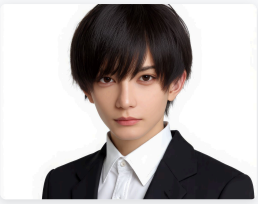

↑ The reference face. Every scene had to preserve this identity.

BEFORE — NO LORA

A different person generated each time

AFTER — WITH LORA

Consistent face maintained

How I solved it

Whisk and ChatGPT image referencing drifted after a few images. Solution: generate 31 expressions with Gemini, then LoRA-train on Replicate's Flux1.dev. Five runs to find optimal parameters.

Full parameter logs from all 5 LoRA training runs and decision rationale are in the paid course. Details will be sent to the email address you registered.

Challenge 2: LoRA Only Produced "Clean" Images

What happened

LoRA preserved character consistency, but horror effects (face covered in dark liquid) became impossible. Even with "black paint" in the prompt, the model defaulted to clean outputs.

BEFORE — LORA ONLY

Dirt/damage effects wouldn't render

AFTER — CHATGPT IMAGE GEN

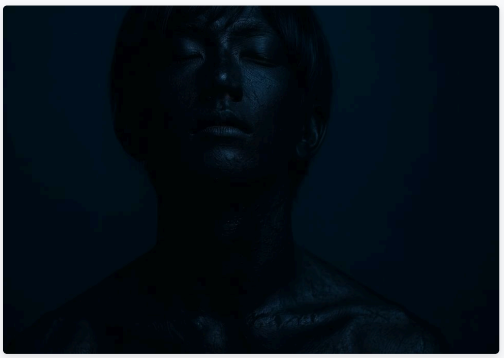

Visceral texture achieved

How I solved it

Root cause: training on 25 "clean" images biased the model. I used ChatGPT image gen to create dark-paint images directly, using terms like "organic mass of meat," "charred," "coagulated" for visceral texture.

Test results, full prompts, and the LoRA v1–v5 evolution log are in the paid course. Details will be sent to the email address you registered.

Challenge 3: Every Method Before LoRA Failed

What happened

Before LoRA, I tried character consistency via Google Whisk + Gemini. Hair and facial features varied every time — far below film quality.

BEFORE — WHISK + GEMINI

Face unstable across generations

AFTER — WITH LORA

Consistent character maintained

Methods that failed

Technical prompts ("85mm lens," etc.) made the AI focus on camera specs, reducing facial accuracy. Pasting images into Gemini chat beat Whisk but added watermarks. LoRA was the only viable long-term solution.

Five Whisk prompt variations and their test results are in the paid course. Details will be sent to the email address you registered.

Beyond these three, I overcame four more: content policy restrictions, dark-image upscale artifacts, video motion limits, and others. All seven are documented in the paid course.

The Short-to-Long Challenge

AI video tools output max 10 seconds per clip. Building 39 minutes (2,340 sec) under this limit is the core challenge every AI filmmaker faces.

The Core Problem

Runway Gen-4 caps at 10 seconds. Its extend feature causes abrupt color/motion shifts. Naive concatenation breaks dialogue-visual sync.

Solution: Forward-Reverse Lip-Sync

Alternate a 10-sec "forward" clip with its reverse. For mostly-static indoor footage, the reverse is nearly imperceptible.

Step	Action
1	Generate a silent 10-sec "forward" clip in Runway
2	Create a reversed version using the reverse function
3	Produce dialogue + silence = 10 sec audio via Hailuo Audio
4	Trim audio to exactly 10.0 sec in Clipchamp
5	Upload 10-sec audio to Runway → lip-sync video
6	Interleave forward → reverse in Premiere Pro

Result: 10 sec × 2 = 20 seconds per cycle. Repeating this, I achieved long takes of ~2 minutes with visual continuity and lip-sync intact.

Managing 39 Minutes as 17 Sequences

I split the 39-minute film into 17 sequences, each managed independently and assigned the optimal method by type — dialogue, horror, narration, etc.

Total Sequences	**17**
Total Runtime	**39 min (~2,340 sec)**
Estimated 10-sec Clips	**234+**
Editing Software	**Adobe Premiere Pro**

Pre-Calculation Was Essential

Dialogue length, pauses, and emotional pacing had to be calculated precisely before production. I used SRT files for second-level timing design.

Hailuo Audio's <#2.9#> syntax inserts 2.9 sec of silence. "Silence + dialogue + silence = 10.0 sec" uploaded to Runway triggers lip movement only during speech.

Limitations exist: lighting changes make reversal slightly visible, and Runway doesn't always output exactly 10 sec. Not perfect, but the most practical long-form AI filmmaking method available.

SRT timecode examples, Hailuo silence syntax guide, and all 17 sequence templates are in the paid course. Details will be sent to the email address you registered.

Next Steps

Thank you for reading this far.
I've covered the overview and key challenges,
but prompts, parameters, and templates remain.

What's in the Paid Course

Runway Gen-4 Scene Prompts — ~62 total
Full English prompts + design rationale for every scene

LoRA Training Parameter Log (v1–v5)
What changed each run and why

BGM Prompts — All 12 Tracks
SUNO / Udio / ElevenLabs comparison included

Hailuo Audio Voice Settings — 7 Characters
Pitch, volume, and emotion configuration

ElevenLabs SFX Prompts — All 24
Selection rationale + failure analysis of 7 rejected attempts

Before/After Images — Full Set + Technical Notes
Includes 4 cases not covered in this document

Long-Form Management Templates
SRT timecode calculator + forward-reverse design examples

Details will be sent to
the email address you registered.

Koki Yoshikawa — Kyoto University, Faculty of Science
HALUCINATION (2025) — Director, Writer, Producer